

# term-project-underwriter

October 29, 2024

## 1 Case Study: “Underwriter for a Day”

Tejas Bajaj 11308438

Haotong Sun 11337172

Yilun Yang 11313432

October 28, 2024

1.1 1.Risk Group Assignment Algorithm: Develop an algorithm for categorizing new applicants into specific risk groups based on their individual characteristics.

```
[ ]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[ ]: # Load your dataset
claim_data = pd.read_csv('claim_data_group5_2024.csv')
```

```
[ ]: claim_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100000 entries, 0 to 99999
Data columns (total 13 columns):
#   Column          Non-Null Count  Dtype
---  -
0   IDpol            100000 non-null float64
1   ClaimNb          100000 non-null float64
2   Exposure         100000 non-null float64
3   Area             100000 non-null object
4   VehPower         100000 non-null float64
5   VehAge           100000 non-null float64
6   DrivAge          100000 non-null float64
7   BonusMalus       100000 non-null float64
8   VehBrand         100000 non-null object
9   VehGas           100000 non-null object
```

```

10 Density      100000 non-null float64
11 Region       100000 non-null object
12 ClaimAmount  100000 non-null float64
dtypes: float64(9), object(4)
memory usage: 9.9+ MB

```

```
[ ]: claim_data
```

	IDpol	ClaimNb	Exposure	Area	VehPower	VehAge	DrivAge	\
0	2271893.0	0.0	0.83	E	5.0	17.0	53.0	
1	1111864.0	0.0	0.24	E	5.0	2.0	27.0	
2	72908.0	0.0	0.50	E	7.0	11.0	67.0	
3	2283027.0	0.0	0.08	B	5.0	8.0	28.0	
4	1123838.0	0.0	0.03	A	11.0	1.0	38.0	
...	...	...	...	...	...	...	...	
99995	70445.0	0.0	1.00	C	5.0	11.0	37.0	
99996	4163362.0	0.0	0.22	E	6.0	13.0	58.0	
99997	2081912.0	0.0	1.00	E	5.0	1.0	49.0	
99998	2012998.0	0.0	0.71	D	9.0	9.0	36.0	
99999	3087666.0	0.0	0.53	C	9.0	14.0	35.0	

	BonusMalus	VehBrand	VehGas	Density	Region	ClaimAmount
0	64.0	B2	Diesel	3317.0	R93	0.0
1	64.0	B3	Diesel	2740.0	R22	0.0
2	50.0	B3	Regular	4762.0	R93	0.0
3	60.0	B1	Diesel	64.0	R91	0.0
4	50.0	B2	Regular	16.0	R24	0.0
...	...	...	...	...	...	...
99995	56.0	B2	Diesel	317.0	R82	0.0
99996	50.0	B1	Diesel	4762.0	R93	0.0
99997	50.0	B2	Diesel	4998.0	R11	0.0
99998	54.0	B1	Regular	1541.0	R91	0.0
99999	51.0	B3	Regular	161.0	R31	0.0

[100000 rows x 13 columns]

```
[ ]: # Descriptive Statistics for Numerical Columns
descriptive_stats = claim_data.describe()
descriptive_stats
```

	IDpol	ClaimNb	Exposure	VehPower	\
count	1.000000e+05	100000.000000	100000.000000	100000.000000	
mean	2.617735e+06	0.039020	0.528057	6.460230	
std	1.643394e+06	0.206296	0.364232	2.055641	
min	1.500000e+01	0.000000	0.002732	4.000000	
25%	1.156127e+06	0.000000	0.170000	5.000000	
50%	2.271008e+06	0.000000	0.490000	6.000000	
75%	4.044791e+06	0.000000	0.990000	7.000000	

max	6.114324e+06	4.000000	1.000000	15.000000
-----	--------------	----------	----------	-----------

	VehAge	DrivAge	BonusMalus	Density \
count	100000.000000	100000.000000	100000.000000	100000.000000
mean	6.992550	45.483040	59.822980	1800.69569
std	5.637297	14.154698	15.652541	3955.08311
min	0.000000	18.000000	50.000000	2.000000
25%	2.000000	34.000000	50.000000	94.000000
50%	6.000000	44.000000	50.000000	399.000000
75%	11.000000	55.000000	65.000000	1658.000000
max	100.000000	99.000000	230.000000	27000.000000

	ClaimAmount
count	100000.000000
mean	76.599887
std	1531.841302
min	0.000000
25%	0.000000
50%	0.000000
75%	0.000000
max	200000.000000

```
[ ]: # Analyzing unique values for categorical columns
```

```
categorical_overview = {
    "Area": claim_data["Area"].nunique(),
    "VehBrand": claim_data["VehBrand"].nunique(),
    "VehGas": claim_data["VehGas"].unique(),
    "Region": claim_data["Region"].nunique()
}
categorical_overview
```

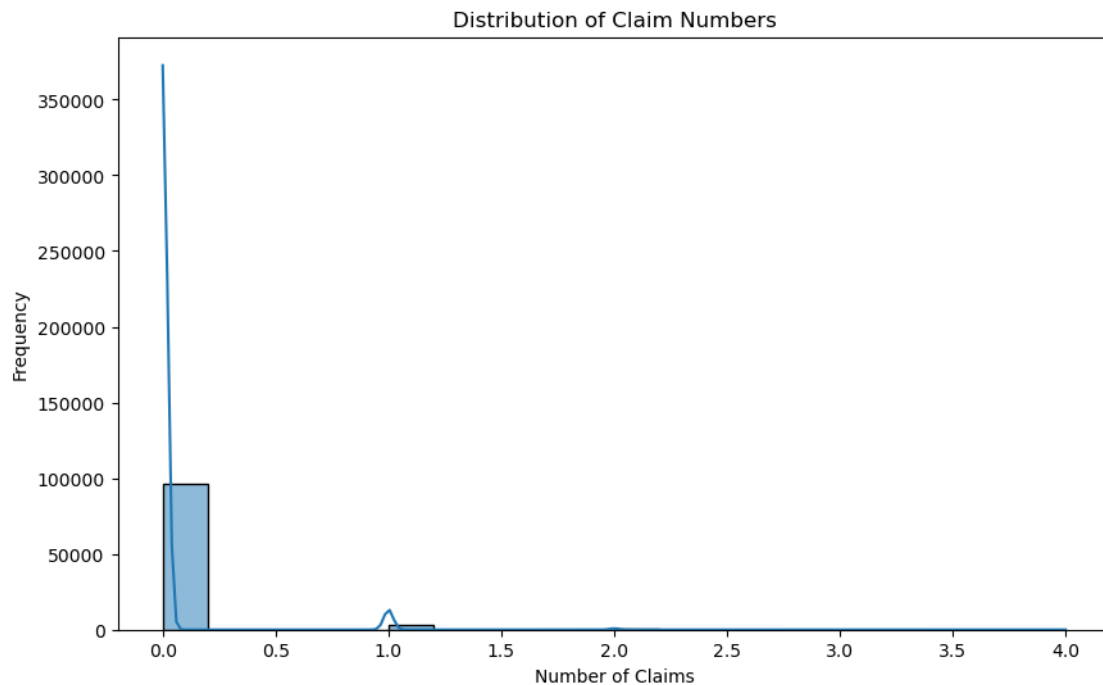
```
{'Area': 6,
 'VehBrand': 11,
 'VehGas': array(['Diesel', 'Regular'], dtype=object),
 'Region': 22}
```

```
[ ]: # Distribution of Individual Claims (ClaimNb)
```

```
claim_distribution = claim_data['ClaimNb'].value_counts()
claim_distribution
```

```
0.0    96335
1.0     3441
2.0      214
3.0        7
4.0         3
Name: ClaimNb, dtype: int64
```


```
[ ]: plt.figure(figsize=(10, 6))
sns.histplot(claim_data['ClaimNb'], bins=20, kde=True).set_title('Distribution_
of Claim Numbers')
plt.xlabel('Number of Claims')
plt.ylabel('Frequency')
plt.show()
```



```
[ ]: # Distribution of Claim Amounts (ClaimAmount)
claim_amount_distribution = claim_data['ClaimAmount'].value_counts()
claim_amount_distribution
```


```
0.00      96335
1204.00     649
1128.12     398
1172.00     282
1128.00      96
...
2487.55      1
1307.64      1
1858.81      1
741.77       1
1117.64      1
Name: ClaimAmount, Length: 1938, dtype: int64
```

```
[ ]: plt.figure(figsize=(10, 6))
sns.histplot(claim_data['ClaimAmount'], bins=50, kde=True).
    ↪ set_title('Distribution of Claim Amounts')
plt.xlabel('Claim Amount')
plt.ylabel('Frequency')
plt.yscale('log') # Using log scale due to large range of values
plt.show()
```



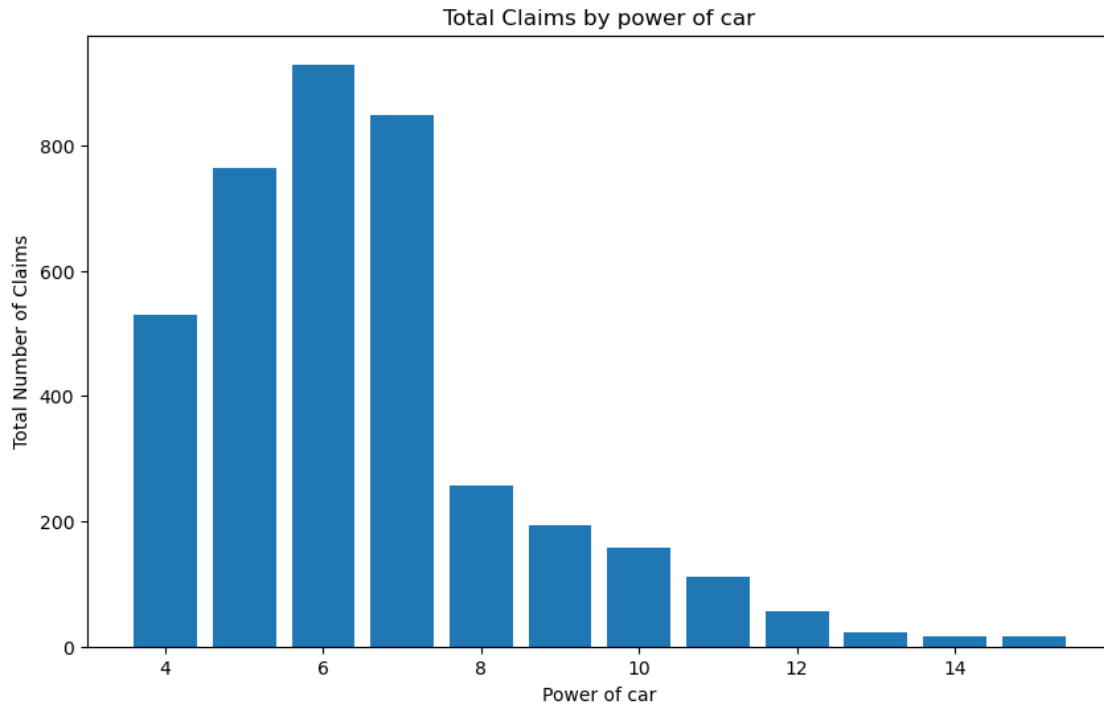
```
[ ]: # Correlation matrix for numerical features
correlation_matrix = claim_data[['ClaimNb', 'ClaimAmount', 'IDpol', 'Exposure', 'VehPower', 'VehAge', 'DrivAge', 'BonusMalus', 'Density']].corr()

plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm')
plt.title('Correlation Matrix')
plt.show()
```



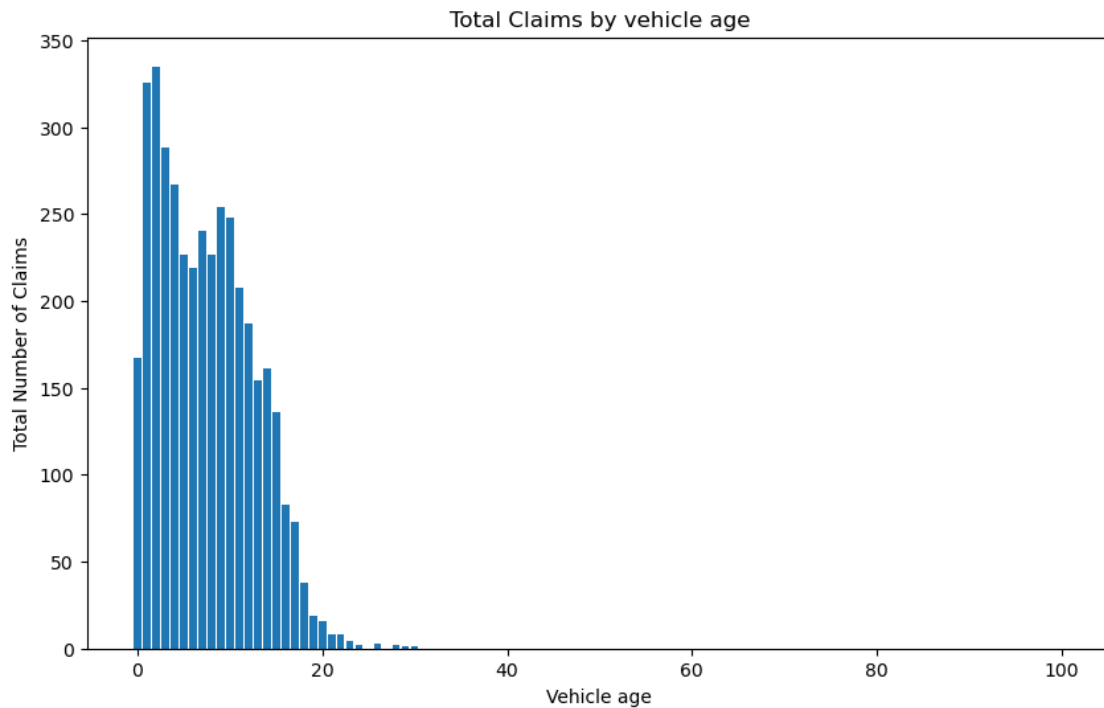
```
[ ]: # Total Claims by power of car (VehPower)
claims_by_power = claim_data.groupby('VehPower')['ClaimNb'].sum().reset_index()

# Plot Claims by power of car
plt.figure(figsize=(10, 6))
plt.bar(claims_by_power['VehPower'], claims_by_power['ClaimNb'])
plt.title('Total Claims by power of car')
plt.xlabel('Power of car')
plt.ylabel('Total Number of Claims')
plt.show()
```



```
[ ]: # Total Claims by vehicle age (VehAge)
claims_by_vehicle_age = claim_data.groupby('VehAge')['ClaimNb'].sum().
    ↪reset_index()


# Plot Claims by vehicle age
plt.figure(figsize=(10, 6))
plt.bar(claims_by_vehicle_age['VehAge'], claims_by_vehicle_age['ClaimNb'])
plt.title('Total Claims by vehicle age')
plt.xlabel('Vehicle age')
plt.ylabel('Total Number of Claims')
plt.show()
```



```
[ ]: # Total Claims by Driver's Age (DrivAge)
claims_by_age = claim_data.groupby('DrivAge')['ClaimNb'].sum().reset_index()

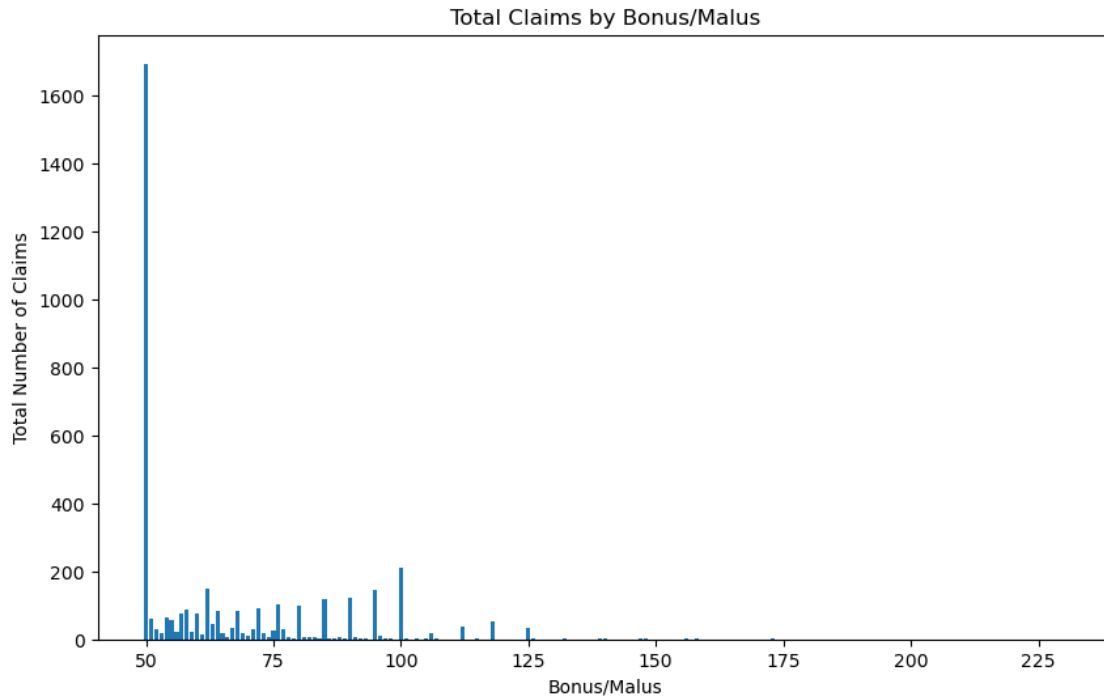
# Plot Claims by Driver's Age
plt.figure(figsize=(10, 6))
plt.bar(claims_by_age['DrivAge'], claims_by_age['ClaimNb'])
plt.title('Total Claims by Driver Age')
plt.xlabel('Driver Age')
plt.ylabel('Total Number of Claims')
plt.show()
```





```
[ ]: # Total Claims by BonusMalus
claims_by_bonusmalus = claim_data.groupby('BonusMalus')['ClaimNb'].sum().
    ↪reset_index()

# Plot Claims by BonusMalus
plt.figure(figsize=(10, 6))
plt.bar(claims_by_bonusmalus['BonusMalus'], claims_by_bonusmalus['ClaimNb'])
plt.title('Total Claims by Bonus/Malus')
plt.xlabel('Bonus/Malus')
plt.ylabel('Total Number of Claims')
plt.show()
```




```
[ ]: # Total Claims by car fuel (VehGas)
claims_by_vehgas = claim_data.groupby('VehGas')['ClaimNb'].sum().reset_index()

# Plotting the pie chart
plt.figure(figsize=(8, 8))
plt.pie(
    claims_by_vehgas['ClaimNb'],
    labels=claims_by_vehgas['VehGas'],
    autopct=lambda pct: f"{pct:.1f}% ({int(pct * claims_by_vehgas['ClaimNb'].
    ↪sum() / 100)})",
    startangle=90
)

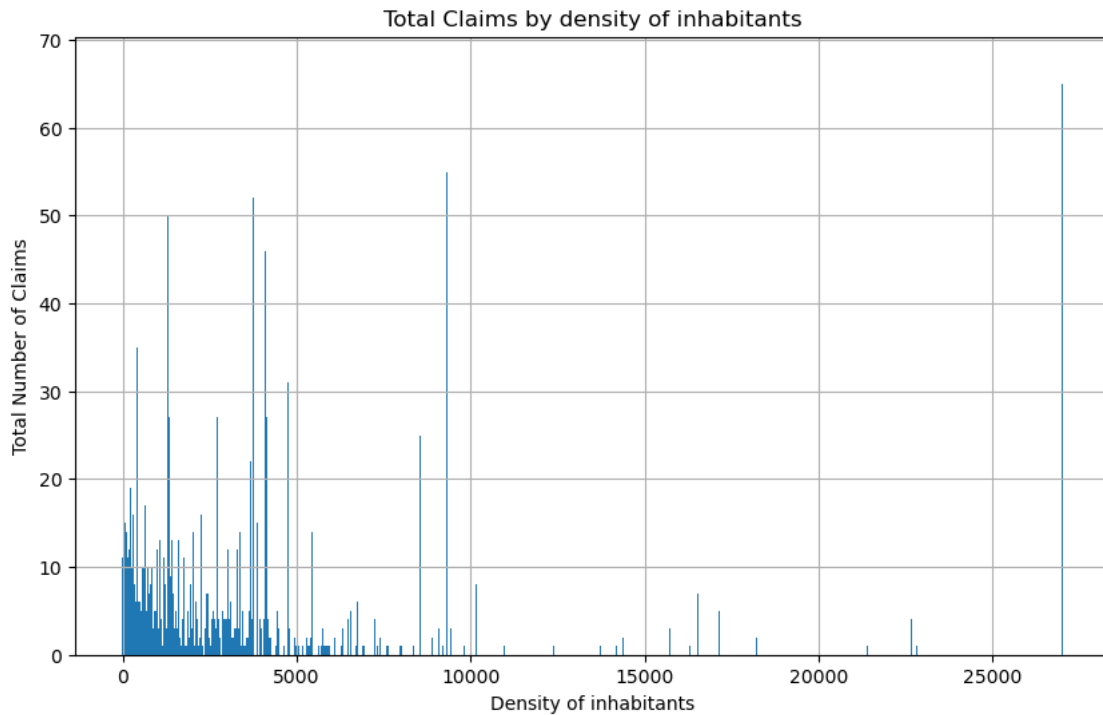
# Adding title
plt.title('Total Claims by Car Fuel')
plt.show()
```

Total Claims by Car Fuel



```
[ ]: # Total Claims by density of inhabitants (Density)
claims_by_density = claim_data.groupby('Density')['ClaimNb'].sum().reset_index()

# Plot Claims by density of inhabitants
plt.figure(figsize=(10, 6))
plt.bar(claims_by_density['Density'], claims_by_density['ClaimNb'], width=30)
plt.title('Total Claims by density of inhabitants')
plt.xlabel('Density of inhabitants')
plt.ylabel('Total Number of Claims')
plt.grid(True)
plt.show()
```



## 1.2 Risk Group Assignment Algorithm

```
[ ]: # Define scoring functions for each characteristic based on the specified rules
```

```
def score_vehicle_power(veh_power):
    if veh_power <= 5:
        return 1 # Low risk
    elif veh_power <= 8:
        return 2 # Moderate risk
    else:
        return 3 # High risk

def score_vehicle_age(veh_age):
    if veh_age <= 5:
        return 1 # Low risk
    elif veh_age <= 10:
        return 2 # Moderate risk
    else:
        return 3 # High risk

def score_driver_age(driver_age):
    if driver_age <= 24:
        return 3 # High risk
```

```

elif driver_age <= 34:
    return 2 # Moderate risk
elif driver_age <= 49:
    return 1 # Low risk
elif driver_age <= 69:
    return 2 # Moderate risk
else:
    return 3 # High risk

def score_bonus_malus(bonus_malus):
    if bonus_malus < 100:
        return 1 # Low risk
    elif bonus_malus <= 200:
        return 2 # Moderate risk
    else:
        return 3 # High risk

def score_density(density):
    if density < 1000:
        return 1 # Low risk
    elif density <= 5000:
        return 2 # Moderate risk
    else:
        return 3 # High risk

def score_fuel_type(fuel_type):
    return 2 if fuel_type == 'Diesel' else 1 # Diesel slightly higher risk
    ↪ than Regular

# Apply the scoring functions to each row in the dataset
claim_data['RiskScore'] = (
    claim_data['VehPower'].apply(score_vehicle_power) +
    claim_data['VehAge'].apply(score_vehicle_age) +
    claim_data['DrivAge'].apply(score_driver_age) +
    claim_data['BonusMalus'].apply(score_bonus_malus) +
    claim_data['Density'].apply(score_density) +
    claim_data['VehGas'].apply(score_fuel_type)
)

# Define thresholds for risk categories based on the total score
def assign_risk_category(score):
    if score <= 7:
        return 'Low Risk'
    elif score <= 12:
        return 'Moderate Risk'
    else:
        return 'High Risk'

```

```
# Assign risk category to each policy based on the total score
claim_data['RiskCategory'] = claim_data['RiskScore'].apply(assign_risk_category)

# Summarize the number of policies in each risk category
risk_category_summary = claim_data['RiskCategory'].value_counts()
risk_category_summary
```

```
Moderate Risk    88620
Low Risk         10118
High Risk        1262
Name: RiskCategory, dtype: int64
```

### 1.2.1 Actuarial Criteria

**Accuracy:** Each characteristic used in the risk categorization system, including vehicle power, vehicle age, driver age, BonusMalus score, population density, and fuel type, has a direct, observable impact on expected costs and losses. For instance, younger drivers typically contribute to higher claim frequencies and severities, making them reliable indicators of risk. **Homogeneity:** By categorizing policies based on clear, measurable characteristics, the system ensures that individuals within the same category share similar risk profiles, reducing the variability of expected claims. The BonusMalus system, in particular, distinguishes drivers based on past claims experience, making each risk group more internally consistent. **Credibility:** The scoring and categorization process was applied to a dataset with a substantial number of policies, especially within the “Moderate Risk” group, thus ensuring the statistical credibility of each category. Each risk group has a sufficient sample size, providing a reliable basis for actuarial assumptions and future predictions. **Reliability:** The criteria used, such as vehicle and driver age, historical claim data, and location-based density, are stable over time. These factors are expected to maintain consistent predictive power, contributing to the stability of the system’s risk predictions.

### 1.2.2 Operational Criteria

**Objectivity:** The risk categories are mutually exclusive and exhaustive, ensuring that each applicant fits into exactly one risk class based on their observable characteristics. The algorithm’s decisions are data-driven, reducing subjective judgment. **Low Costs and Ease of Observation:** Each characteristic—such as driver age, BonusMalus score, and vehicle type—is easily observed, typically self-reported, and already part of standard data collection practices, minimizing costs and administrative burden. **Resistance to Manipulation:** Most factors used in scoring, like age, vehicle power, and claim history, cannot be easily manipulated by the applicant. This maintains the integrity of risk categories and reduces incentives for applicants to misreport information. **Intuitive and Explainable:** The factors used in scoring align with common-sense perceptions of risk. For example, people understand why younger drivers might present higher risks, making it easier to explain premium decisions. **Few Disconnects Between Groups:** The categorization system uses smooth transitions in scoring ranges, with clear thresholds, avoiding sharp or arbitrary disconnects between similar applicants in different risk categories.

### 1.2.3 Social Criteria

**Privacy:** The system avoids sensitive data, such as credit scores and personal health information, thereby respecting applicant privacy. It relies on observable characteristics directly relevant to driving risk, avoiding unnecessary or intrusive data collection. **Causality vs. Correlation:** The characteristics used in categorizing risk, such as driver age, vehicle power, and past claims, have a causal connection to driving risk rather than mere correlation. For example, higher vehicle power often results in more severe accidents, which justifies its use as a factor. **Control:** The use of clear risk factors like BonusMalus scores and vehicle attributes allows SafeRoads to better control risks by understanding and potentially mitigating risk levels through incentives for safe driving (e.g., improving BonusMalus scores). **Affordability and Availability:** This risk-based pricing ensures that low-risk customers can access affordable premiums, reducing cross-subsidization where low-risk individuals subsidize high-risk individuals' claims. By fairly pricing risk, the system improves affordability and availability for most insured individuals.

### 1.2.4 Legal Criteria

**Permissibility:** The characteristics used in the categorization, including age, vehicle attributes, and claims history, are all legally recognized and widely used in the insurance industry. They align with regulatory norms for permissible risk factors. **Avoiding Undue Discrimination:** The system is designed to avoid unfair discrimination, using only statistically justifiable and commonly accepted risk factors. Each characteristic has been validated for its predictive accuracy, ensuring that differences in treatment are based on actual risk rather than bias.

## 1.3 Conclusion

According to the data information provided, this risk categorization system is comprehensive, aligning well with actuarial precision and operational efficiency, while respecting social and legal norms. By leveraging data-driven, observable, and fair criteria, SafeRoads can ensure accurate premium setting, financial stability, and customer trust. The system thus meets industry standards and provides a sound foundation for pricing decisions for the upcoming year. But it is undeniable that there is still room for improvement in this risk level classification.

## 1.4 Predicting Total Losses: Based on historical claims data and uElizing probability distributions, you need to model the total losses expected for the upcoming year.

```
[ ]: import pandas as pd

claim_data = pd.read_csv('claim_data_group5_2024.csv')
print("Missing values:\n", claim_data.isnull().sum())
claim_data = claim_data.dropna()
claim_data = claim_data[(claim_data['ClaimNb'] >= 0) &
    ↪(claim_data['ClaimAmount'] >= 0)]
print(claim_data[['ClaimNb', 'ClaimAmount']].describe())
```

```
Missing values:
  IDpol      0
ClaimNb     0
```

Exposure	0
Area	0
VehPower	0
VehAge	0
DrivAge	0
BonusMalus	0
VehBrand	0
VehGas	0
Density	0
Region	0
ClaimAmount	0
dtype: int64	

	ClaimNb	ClaimAmount
count	100000.000000	100000.000000
mean	0.039020	76.599887
std	0.206296	1531.841302
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	0.000000
75%	0.000000	0.000000
max	4.000000	200000.000000

#### 1.4.1 Fit Frequency Distribution

```
[ ]: import pandas as pd
import numpy as np
import scipy.stats as stats
import matplotlib.pyplot as plt

claim_data = pd.read_csv('claim_data_group5_2024.csv')
claim_data = claim_data.dropna()
claim_data = claim_data[(claim_data['ClaimNb'] >= 0) &
    ↳ (claim_data['ClaimAmount'] > 0)] # Ensure ClaimAmount > 0
claim_frequency = claim_data['ClaimNb']
mean_claim_freq = claim_frequency.mean()
var_claim_freq = claim_frequency.var()
if var_claim_freq > mean_claim_freq:
    alpha_est = mean_claim_freq**2 / (var_claim_freq - mean_claim_freq)
    p_est = mean_claim_freq / var_claim_freq
else:
    alpha_est = mean_claim_freq
    p_est = 1.0

print("Estimated Negative Binomial parameters:")
print("Alpha (dispersion parameter):", alpha_est)
print("P (probability of success):", p_est)
```



```

# Plot
plt.hist(claim_frequency, bins=range(0, int(claim_frequency.max()) + 1),
         density=True, alpha=0.6, color='g', label='Empirical')

x = np.arange(0, int(claim_frequency.max()) + 1)
if var_claim_freq > mean_claim_freq:
    fitted_nb = stats.nbinom.pmf(x, alpha_est, p_est)
else:
    fitted_nb = stats.poisson.pmf(x, alpha_est)
plt.plot(x, fitted_nb, 'r-', label='Fitted Distribution')

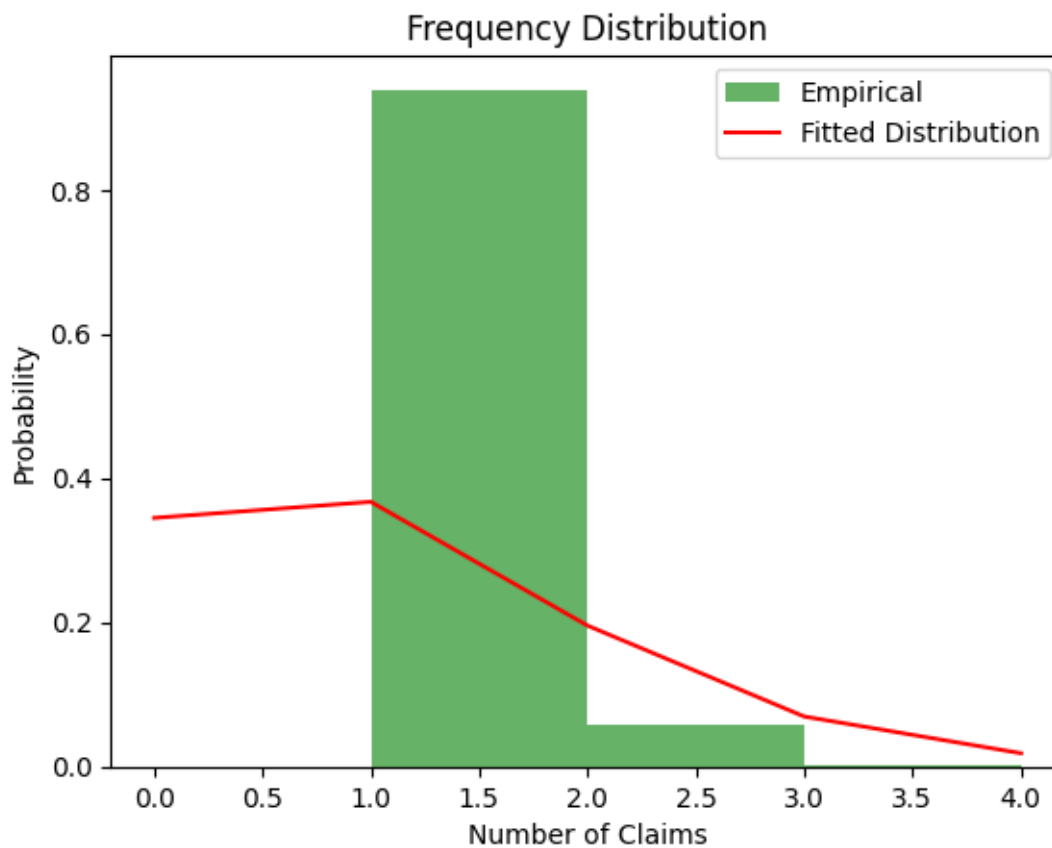
plt.title('Frequency Distribution')
plt.xlabel('Number of Claims')
plt.ylabel('Probability')
plt.legend()
plt.show()

```

Estimated Negative Binomial parameters:

Alpha (dispersion parameter): 1.0646657571623466

P (probability of success): 1.0



Here we estimated the dispersion parameter Alpha, which is 1.0646657571623466, it shows there's a high tendency of dispersion but the alpha is around 1 shows the dispersion is not extremely high. And the probability of success is 1.0 means it follows a poisson distribution.

### Simulated loss Gamma

According to the simulation, the average total expected loss in a year would be 2000.56 with a standard deviation of 3676.91, indicating a high degree of volatility in the total loss.

```
[ ]: # Severity distribution
claim_severity = claim_data['ClaimAmount']
alpha_sev, loc_sev, beta_sev = stats.gamma.fit(claim_severity, floc=0)

print("Estimated Gamma parameters for severity distribution:")
print("Alpha:", alpha_sev)
print("Loc:", loc_sev)
print("Beta:", beta_sev)

# Simulate total losses for the next year
num_simulations = 10000
total_losses = []

for _ in range(num_simulations):
    if var_claim_freq > mean_claim_freq:
        simulated_claims = np.random.negative_binomial(alpha_est, p_est)
    else:
        simulated_claims = np.random.poisson(alpha_est)

    simulated_severity = np.random.gamma(alpha_sev, beta_sev, simulated_claims)

    # total loss
    total_loss = simulated_severity.sum()
    total_losses.append(total_loss)

# Convert to numpy array
total_losses = np.array(total_losses)

# Plot
plt.hist(total_losses, bins=50, density=True, alpha=0.6, color='g')
plt.title('Simulated Total Losses for Next Year')
plt.xlabel('Total Loss')
plt.ylabel('Frequency')
plt.show()
print("Simulated Total Losses - Mean:", total_losses.mean())
print("Simulated Total Losses - Std Dev:", total_losses.std())
ks_stat, p_value = stats.kstest(claim_severity, 'gamma', args=(alpha_sev, loc_sev, beta_sev))
```


```
print("KS Test for Gamma Distribution - Statistic:", ks_stat, "P-value:",  
      ↪p_value)
```

Estimated Gamma parameters for severity distribution:

Alpha: 0.77598083307877

Loc: 0

Beta: 2693.4142535413143



Simulated Total Losses - Mean: 2195.6363366800383

Simulated Total Losses - Std Dev: 3176.261629979459

KS Test for Gamma Distribution - Statistic: 0.24645803383389342 P-value:  
1.5089949709438602e-196

### Simulated loss (log-normal) and check the best fit

```
[ ]:
```

```
[ ]: # Estimate severity distribution (Log-Normal)
claim_severity = claim_data['ClaimAmount']
shape, loc, scale = stats.lognorm.fit(claim_severity, floc=0)
```

```

# Print estimated parameters
print("Estimated Log-Normal parameters for severity distribution:")
print("Shape:", shape)
print("Loc:", loc)
print("Scale:", scale)

# Simulate total losses for the next year
num_simulations = 10000
total_losses = []

for _ in range(num_simulations):

    if var_claim_freq > mean_claim_freq:
        simulated_claims = np.random.negative_binomial(alpha_est, p_est)
    else:
        simulated_claims = np.random.poisson(alpha_est)

    simulated_severity = np.random.lognormal(np.log(scale), shape,
    ↪ simulated_claims)
    total_loss = simulated_severity.sum()
    total_losses.append(total_loss)

# Convert to numpy array
total_losses = np.array(total_losses)

# Plot
plt.hist(total_losses, bins=50, density=True, alpha=0.6, color='g')
plt.title('Simulated Total Losses for Next Year')
plt.xlabel('Total Loss')
plt.ylabel('Frequency')
plt.show()

# Print basic stats for simulated total losses
print("Simulated Total Losses - Mean:", total_losses.mean())
print("Simulated Total Losses - Std Dev:", total_losses.std())

# Goodness-of-fit test for Log-Normal distribution
ks_stat_lognorm, p_value_lognorm = stats.kstest(claim_severity, 'lognorm',
    ↪ args=(shape, loc, scale))
print("KS Test for Log-Normal Distribution - Statistic:", ks_stat_lognorm,
    ↪ "P-value:", p_value_lognorm)

# Compare
alpha_sev, loc_sev, beta_sev = stats.gamma.fit(claim_severity, floc=0)
ks_stat_gamma, p_value_gamma = stats.kstest(claim_severity, 'gamma',
    ↪ args=(alpha_sev, loc_sev, beta_sev))

```

```

print("KS Test for Gamma Distribution - Statistic:", ks_stat_gamma, "P-value:", p_value_gamma)

# Choose the best fit based on p-value
if p_value_lognorm > p_value_gamma:
    best_fit = 'lognorm'
    best_params = (shape, loc, scale)
else:
    best_fit = 'gamma'
    best_params = (alpha_sev, loc_sev, beta_sev)

print(f"Best fit distribution: {best_fit} with parameters: {best_params}")

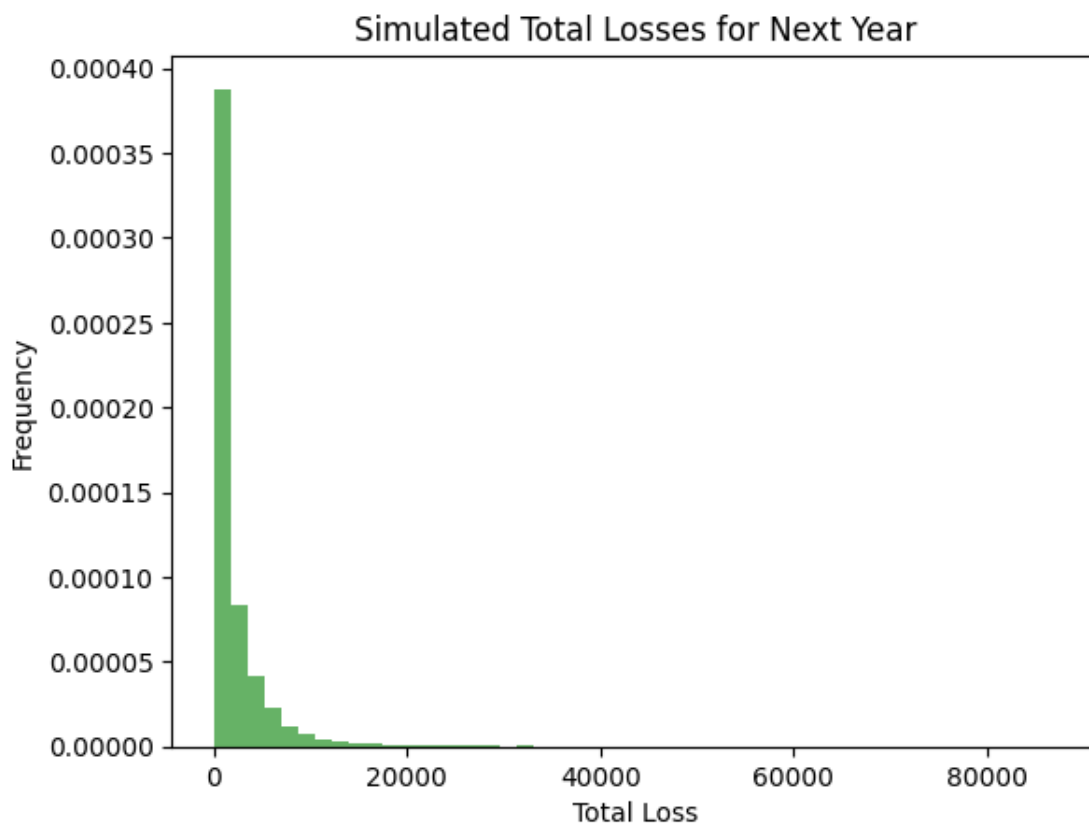
```

Estimated Log-Normal parameters for severity distribution:

Shape: 1.1596764669034225

Loc: 0

Scale: 969.7229460522528



Simulated Total Losses - Mean: 2000.5555337452777

Simulated Total Losses - Std Dev: 3676.914082236124

KS Test for Log-Normal Distribution - Statistic: 0.2307185397172155 P-value:

5.315203177357836e-172

KS Test for Gamma Distribution - Statistic: 0.24645803383389342 P-value:  
1.5089949709438602e-196

Best fit distribution: lognorm with parameters: (1.1596764669034225, 0,  
969.7229460522528)

According to the simulation of log-normal distribution, the average total expected loss in a year would be 2000.56 with a standard deviation of 3676.91 while the gamma distribution is 2195.63 and 3176.26, indicating a high degree of volatility in the total loss. From the results of the Kolmogorov-Smirnov test, the lognormal distribution performs better than the gamma distribution in fitting the data (higher p-value). Lognormal distribution has a statistic value of 0.2307 and gamma distribution has a statistic value of 0.2465, both having lower P-value, which indicates that despite the fact that both are imperfect, the lognormal distribution has a relatively better fit .

### 1.5 3. Premium Determination: Assign appropriate premiums to each risk group, ensuring that the likelihood of not being able to cover claims next year does not exceed 0.5%.

```
[1]: import pandas as pd
import numpy as np
from scipy.stats import norm
```

```
[3]: claim_data = pd.read_csv('claim_data_group5_2024.csv')
claim_data.head()
```

```
[3]:
```

	IDpol	ClaimNb	Exposure	Area	VehPower	VehAge	DrivAge	BonusMalus	
0	2271893.0	0.0	0.83	E	5.0	17.0	53.0	64.0	\
1	1111864.0	0.0	0.24	E	5.0	2.0	27.0	64.0	
2	72908.0	0.0	0.50	E	7.0	11.0	67.0	50.0	
3	2283027.0	0.0	0.08	B	5.0	8.0	28.0	60.0	
4	1123838.0	0.0	0.03	A	11.0	1.0	38.0	50.0	

	VehBrand	VehGas	Density	Region	ClaimAmount
0	B2	Diesel	3317.0	R93	0.0
1	B3	Diesel	2740.0	R22	0.0
2	B3	Regular	4762.0	R93	0.0
3	B1	Diesel	64.0	R91	0.0
4	B2	Regular	16.0	R24	0.0

```
[4]: # Define driver age categories and vehicle power levels as risk group
↳ identifiers
claim_data['DrivAgeGroup'] = pd.cut(claim_data['DrivAge'], bins=[18, 25, 35,
↳ 50, 65, 100], labels=['18-25', '26-35', '36-50', '51-65', '65+'])
claim_data['RiskGroup'] = claim_data['DrivAgeGroup'].astype(str) + "_" +
↳ claim_data['VehPower'].astype(str)
```

```
[5]: # Calculate claim frequency and severity per risk group
risk_group_stats = claim_data.groupby('RiskGroup').agg(
    frequency=('ClaimNb', 'mean'), # Average claims per policy
    severity=('ClaimAmount', lambda x: x[x > 0].mean()) # Average claim amount
    ↪for policies with claims
).fillna(0) # Replace NaN values with 0 for groups with no claims

risk_group_stats.reset_index(inplace=True)
risk_group_stats.head()
```

```
[5]:
```

	RiskGroup	frequency	severity
0	18-25_10.0	0.036036	718.7925
1	18-25_11.0	0.122449	1235.8700
2	18-25_12.0	0.000000	0.0000
3	18-25_13.0	0.000000	0.0000
4	18-25_14.0	0.187500	746.7100

```
[6]: # Calculate expected loss for each risk group
risk_group_stats['expected_loss'] = risk_group_stats['frequency'] *
    ↪risk_group_stats['severity']
risk_group_stats.head()
```

```
[6]:
```

	RiskGroup	frequency	severity	expected_loss
0	18-25_10.0	0.036036	718.7925	25.902432
1	18-25_11.0	0.122449	1235.8700	151.331020
2	18-25_12.0	0.000000	0.0000	0.000000
3	18-25_13.0	0.000000	0.0000	0.000000
4	18-25_14.0	0.187500	746.7100	140.008125

```
[7]: # Assume a buffer multiplier based on 99.5% confidence using the normal
    ↪distribution (z-score ~2.576)
z_score_99_5 = norm.ppf(0.995)

# Estimate premium as expected loss plus buffer
risk_group_stats['premium'] = risk_group_stats['expected_loss'] * (1 +
    ↪z_score_99_5)
risk_group_stats[['RiskGroup', 'frequency', 'severity', 'expected_loss',
    ↪'premium']]
```

```
[7]:
```

	RiskGroup	frequency	severity	expected_loss	premium
0	18-25_10.0	0.036036	718.7925	25.902432	92.622677
1	18-25_11.0	0.122449	1235.8700	151.331020	541.133897
2	18-25_12.0	0.000000	0.0000	0.000000	0.000000
3	18-25_13.0	0.000000	0.0000	0.000000	0.000000
4	18-25_14.0	0.187500	746.7100	140.008125	500.645156
..	...	...	...	...	...
64	nan_5.0	0.071429	3878.8050	277.057500	990.710327

65	nan_6.0	0.000000	0.0000	0.000000	0.000000
66	nan_7.0	0.130435	2087.5300	272.286522	973.650123
67	nan_8.0	0.000000	0.0000	0.000000	0.000000
68	nan_9.0	0.000000	0.0000	0.000000	0.000000

[69 rows x 5 columns]

```
[8]: # Display the final risk group statistics with calculated premiums
print(risk_group_stats[['RiskGroup', 'frequency', 'severity', 'expected_loss', 'premium']])
```

	RiskGroup	frequency	severity	expected_loss	premium
0	18-25_10.0	0.036036	718.7925	25.902432	92.622677
1	18-25_11.0	0.122449	1235.8700	151.331020	541.133897
2	18-25_12.0	0.000000	0.0000	0.000000	0.000000
3	18-25_13.0	0.000000	0.0000	0.000000	0.000000
4	18-25_14.0	0.187500	746.7100	140.008125	500.645156
..	...	...	...	...	...
64	nan_5.0	0.071429	3878.8050	277.057500	990.710327
65	nan_6.0	0.000000	0.0000	0.000000	0.000000
66	nan_7.0	0.130435	2087.5300	272.286522	973.650123
67	nan_8.0	0.000000	0.0000	0.000000	0.000000
68	nan_9.0	0.000000	0.0000	0.000000	0.000000

[69 rows x 5 columns]